

# Variable Selection

Group 5

December 12, 2015

In order to find the outliers, we will train a model using a section of the dataset that does not include our first 400 observations. We will then fit that model to the dataset we are interested in. This will avoid the 50 outliers from impacting the estimated coefficients in such a way as to mask themselves. We have therefore used observations 401 to 3000 of the synthetic regression data to train a model, which we will test on observations 1 to 400.

We will use a stepwise algorithm in order to find out which of the variables to include in the dataset. The major flaw in this method is that it fails to take into account correlations between the different variables. In order to get around this issue, we create groups of highly correlated variables. Therefore instead of, for instance, making a stepwise choice whether to include variable one, variable two, or variable three, the algorithm chooses whether to include variables one or two *and* three together.

```
file <- '/home/beeb/Documents/Data_Science/Data/Stat/'
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

setwd(file)
synth.reg.test <- read.table('synthetic_regression.txt', header = TRUE, nrows= 400)
synth.reg.train <- read.table('synthetic_regression.txt', header = TRUE, nrows = 3000)
synth.reg.train <- synth.reg.train[401:3000,]
```

This section of code identifies the correlated x's:

```
# This bit of code figures out which of the xs need to be put in groups together
# The first thing we're going to do is see if any of the x's are correlated
cors <- cor(synth.reg.train[2:ncol(synth.reg.train)])
m <- ncol(cors)

# Pick out those that have correlation above 0.5
together <- apply(abs(cors)>0.5, 2, which)
```

We now have a list, 'together', which contains all the variables that should go together. We will scan through this list and see which groups of variables we can add to our current set of variables in order to provide a model with the highest log likelihood.

```

# Initialise a dataframe for adding variables to and a list to show which groups were added
currentmod <- data.frame(t = synth.reg.train$t)
incvars <- list()

# Stepwise!
# Take the first 50 most important variables (we can examine subsets later on)
for(k in 1:50) {
  likelihoods <- sapply(together, function(y) {
    mod <- cbind(currentmod, select(synth.reg.train, one_of(names(y))))
    model <- lm(t ~ ., data = mod)
    return(logLik(model))
  })
  # Choose the variable which will give the best log likelihood
  bestvar <- which(likelihoods == max(likelihoods))
  incvars[[k]] <- bestvar

  # Move the best variables from the list of vars under consideration to the dataframe of selected variables
  together <- together[setdiff(1:length(together), bestvar)]
  currentmod <- cbind(currentmod, select(synth.reg.train, one_of(names(bestvar))))
}

```

We now have a list of the top 50 groups of variables (59 variables in total) that impact on  $t$ . We will now create a series of linear models on the training data, test them on the testing data, and choose a model which provides a high  $R^2$  on the test data.

```

# This bit exists due to the complication of needing to treat certain variables in groups
# That is, highly correlated variables should go together
# 'Steps' will tell us; first take the first three vars; then the next two vars; then
# one by itself ; etc etc.
steps <- cumsum(sapply(incvars, length))
incvars2 <- unlist(incvars)
collect.r2 <- rep(0, length(steps))
num.outliers <- rep(0, length(steps))
j <- 1

for(i in steps) {
  currentmod <- select(synth.reg.train, t, one_of(names(incvars2)[1:i]))
  model <- lm(t ~ ., data = currentmod)
  assign(paste0('model', j), model)
  synth.reg.test$predvals <- predict(model, newdata = synth.reg.test)

  #We collect the R^2 of using the training model on the testing data.
  # Not sure if there's an automatic way to do this.
  # Also mark out the points with low likelihood
  synth.reg.test$residuals <- synth.reg.test$predvals - synth.reg.test$t

  ressumsquare <- sum((synth.reg.test$residuals)**2)
  totsumsquare <- sum((synth.reg.test$t - mean(synth.reg.test$t))**2)
  r.squared <- 1 - (ressumsquare/totsumsquare)
  collect.r2[j] <- r.squared

  # We create a cutoff - a 1.96 standard deviation confidence interval, using the standard deviations
  cutoff <- sd(model$residuals) * 1.96
  synth.reg.test$cutoff <- 0
}

```

```

synth.reg.test$cutoff[abs(synth.reg.test$residuals) > cutoff] <- 1
num.outliers[j] <- sum(synth.reg.test$cutoff)
assign(paste0('synth.reg.test', j), synth.reg.test)

# Now make a graph showing predicted values vs actual values
plottest <- ggplot(data = synth.reg.test, aes(x = predvals, y = t, colour = cutoff)) +
  geom_point() +
  geom_abline(intercept = 0, slope = 1) +
  geom_abline(intercept = cutoff, slope = 1) +
  geom_abline(intercept = -cutoff, slope = 1) +
  ggtitle(paste('first', i, 'vars, R2:', round(r.squared, 3), 'Observations above cutoff:', sum(synth
assign(paste0('plottest', j), plottest)

# And another graph showing the size of the residuals
plottest.res <- ggplot(data = synth.reg.test, aes(y = abs(residuals), x = rownames(synth.reg.test), c
  geom_hline(yintercept = 1.96 * sd(model$residuals), colour = 'deeppink2') +
  ggtitle(paste('first', i, 'vars, R2:', round(r.squared, 3), 'Observations above cutoff:', sum(synth
  xlab('') +
  ylab('residual size')
assign(paste0('plottest.res', j), plottest.res)

j <- j + 1
}

```

We now have a series of graphs which show us how our residuals change depending on the size of our models. We choose the model which gives the highest  $R^2$  when used on the testing dataset. We find it is model 7.

```
which(collect.r2 == max(collect.r2))
```

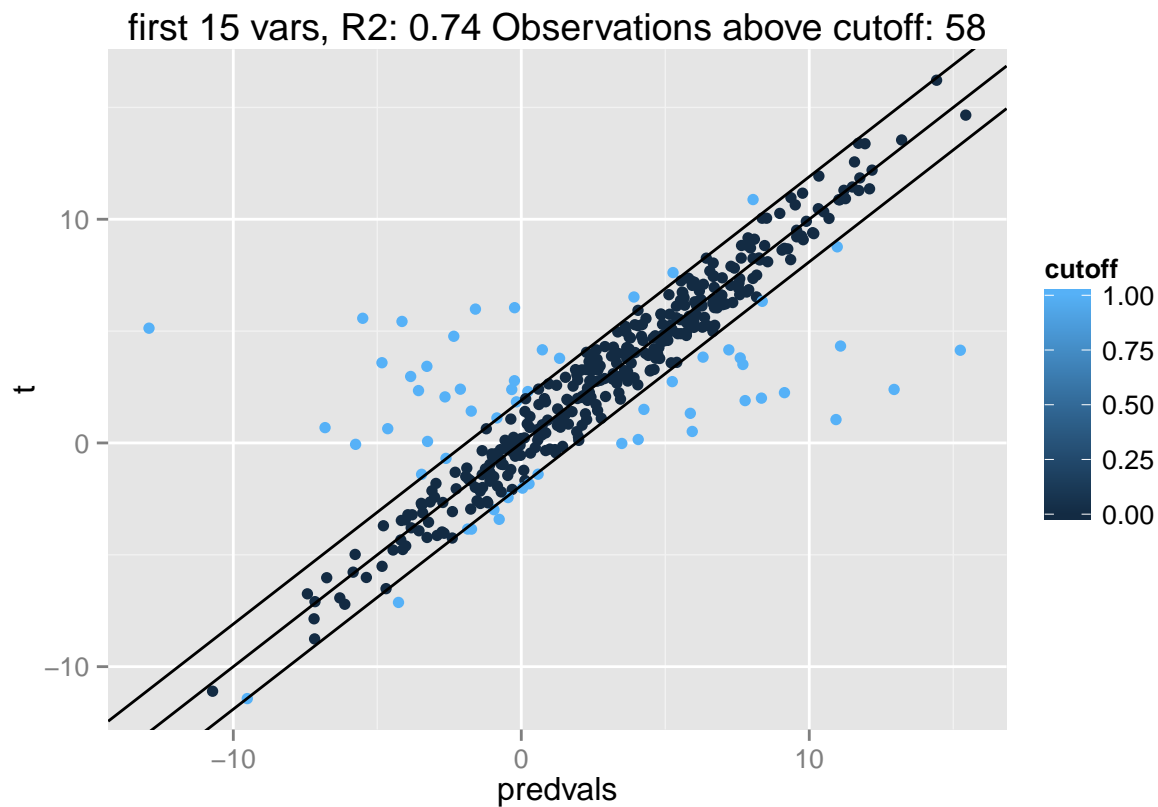
```
## [1] 7
```

```
summary(model7)
```

```
##
## Call:
## lm(formula = t ~ ., data = currentmod)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0955 -0.6501 -0.0103  0.6579  3.4928
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.030107   0.019123 158.457 < 2e-16 ***
## X.1000       2.993732   0.019489 153.612 < 2e-16 ***
## X.5         2.035153   0.040357  50.429 < 2e-16 ***
## X.13       -0.054631   0.032010  -1.707 0.088002 .
## X.14        0.045564   0.031629   1.441 0.149817
## X.17       -0.035560   0.032495  -1.094 0.273912
## X.18       -0.003725   0.031949  -0.117 0.907197
## X.400       1.978385   0.041448  47.732 < 2e-16 ***
## X.15        0.007413   0.031439   0.236 0.813619
```


```
## X.16      0.037925  0.032087  1.182 0.237337
## X.100     0.979436  0.040354 24.271 < 2e-16 ***
## X.1       0.914478  0.040644 22.499 < 2e-16 ***
## X.11      0.043636  0.031538  1.384 0.166605
## X.12      0.054545  0.031117  1.753 0.079742 .
## X.588     0.065936  0.018965  3.477 0.000516 ***
## X.696     -0.053901 0.019022 -2.834 0.004639 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9739 on 2584 degrees of freedom
## Multiple R-squared:  0.9533, Adjusted R-squared:  0.953
## F-statistic: 3514 on 15 and 2584 DF, p-value: < 2.2e-16
```

```
plottest7
```



```
plottest.res7
```

first 15 vars, R2: 0.74 Observations above cutoff: 58



The last step is to create a list of which observations are above the cutoff value. As we see from the graphs above, there is not a clear point separation between outliers and non-outliers. I have therefore included in my list all the observations above the cutoff value, although at the margins the choice of whether to list a certain observation as an outlier or not is somewhat arbitrary. 'Index' captures the row numbers of these observations.

```
outliers <- mutate(synth.reg.test7, index = rownames(synth.reg.test7)) %>%
  filter(cutoff == 1) %>%
  arrange(residuals) %>%
  select(t, predvals, residuals, index)
```

outliers

##	t	predvals	residuals	index
## 1	5.13038686	-12.9272312	-18.057618	225
## 2	5.57352935	-5.5059209	-11.079450	224
## 3	5.43193017	-4.1421851	-9.574115	214
## 4	3.58643325	-4.8384388	-8.424872	216
## 5	5.98378389	-1.5923733	-7.576157	237
## 6	0.68500482	-6.8143459	-7.499351	210
## 7	4.76906714	-2.3442211	-7.113288	241
## 8	2.97237251	-3.8412829	-6.813655	250
## 9	3.42484644	-3.2783785	-6.703225	248
## 10	6.04695065	-0.2328766	-6.279827	233
## 11	2.34123534	-3.5688562	-5.910092	219
## 12	-0.06240902	-5.7536238	-5.691215	236
## 13	0.63800906	-4.6373636	-5.275373	222
## 14	2.05923419	-2.6489700	-4.708204	247
## 15	2.40337699	-2.1155492	-4.518926	218

## 16	4.16469739	0.7310353	-3.433662	202
## 17	0.06785632	-3.2555304	-3.323387	212
## 18	1.42297647	-1.7377345	-3.160711	203
## 19	2.78259043	-0.2373898	-3.019980	255
## 20	10.88127644	8.0478079	-2.833469	27
## 21	2.38861024	-0.3308418	-2.719452	223
## 22	6.52431337	3.9118550	-2.612458	96
## 23	3.78334115	1.3289009	-2.454440	246
## 24	7.61622474	5.2683424	-2.347882	17
## 25	2.29778061	0.2285178	-2.069263	122
## 26	-1.40251745	-3.4618085	-2.059291	101
## 27	1.83774358	-0.1683987	-2.006142	211
## 28	1.11821124	-0.8469641	-1.965175	310
## 29	-0.68954119	-2.6150373	-1.925496	63
## 30	-11.42769535	-9.5142400	1.913455	380
## 31	-2.44093139	-0.4602417	1.980690	319
## 32	-1.39735820	0.5835125	1.980871	200
## 33	-3.84842106	-1.8621365	1.986285	364
## 34	-2.98347051	-0.9494068	2.034064	171
## 35	6.33162247	8.3676256	2.036003	355
## 36	-2.01978256	0.0416753	2.061458	382
## 37	-1.82629680	0.2685642	2.094861	184
## 38	-3.85149810	-1.7288872	2.122611	269
## 39	8.76679008	10.9710138	2.204224	112
## 40	3.83583840	6.3108818	2.475043	204
## 41	2.74229596	5.2428260	2.500530	201
## 42	-3.41430550	-0.7675908	2.646715	328
## 43	1.49979513	4.2577856	2.757990	249
## 44	-7.12906538	-4.2638940	2.865171	254
## 45	4.16193898	7.2066081	3.044669	209
## 46	-0.02311067	3.4902847	3.513395	234
## 47	3.80123003	7.6025581	3.801328	235
## 48	0.15339674	4.0600361	3.906639	232
## 49	3.50907818	7.6947921	4.185714	239
## 50	1.32443465	5.8695333	4.545099	226
## 51	0.51382339	5.9365025	5.422679	231
## 52	1.89304611	7.7726707	5.879625	213
## 53	2.00328029	8.3388120	6.335532	238
## 54	4.33193989	11.0857370	6.753797	245
## 55	2.24823310	9.1385439	6.890311	208
## 56	1.04502780	10.9280584	9.883031	207
## 57	2.39203999	12.9459801	10.553940	229
## 58	4.14485011	15.2448676	11.100018	205